

F-CS-WAE: Factorized Class-Structured Spherical Cauchy Wasserstein Auto-Encoders for Label-Guided Generative Clustering

Anonymous Authors

Anonymous Institution

anonymous@example.com

Abstract

Generative clustering requires a latent space that is simultaneously discriminative, reconstructive, and samplable. These requirements often conflict: a representation can form clean clusters while decoding poor prior samples, or reconstruct images accurately while providing little class structure. We study this tension through F-CS-WAE, a factorized class-structured Wasserstein auto-encoder that separates a hyperspherical semantic latent $z_c \in \mathbb{S}^{d_c-1}$ from a Euclidean style latent $z_s \in \mathbb{R}^{d_s}$. The semantic component is regularized by class-conditional Spherical Cauchy priors on the hypersphere, while the style component is matched to a Gaussian prior using MMD. A four-phase schedule first learns reconstruction, then gradually introduces style matching, auxiliary semantic classification, and semantic prior alignment. On CIFAR-10, F-CS-WAE achieves $80.87 \pm 0.52\%$ clustering accuracy over three seeds with FID 83.04 ± 0.86 , substantially improving over available VAE, WAE-MMD, VaDE, and single-latent baselines under the repository protocol. On MNIST, F-CS-WAE reaches 99.15% clustering accuracy and strong reconstruction quality, but a diagnostic analysis reveals a conditional prior mismatch: the global style aggregate resembles $N(0, I)$, yet style codes remain class-dependent, causing naive independent prior samples to decode poorly. This failure mode is not incidental; it exposes a central design issue in factorized generative clustering. We therefore present F-CS-WAE as both a strong empirical model and a diagnostic framework for studying when semantic/style factorization improves or harms sampling. The current evidence suggests that AAAI-level claims should combine clustering gains with conditional style alignment, rather than relying on clustering metrics alone.

1 Introduction

Auto-encoding generative models promise a compact latent representation from which one can reconstruct inputs, sample new data, and analyze structure. In clustering-oriented settings, however, these goals pull the model in different directions. Reconstruction rewards retaining all information needed to reproduce an image. Clustering rewards removing nuisance variation and preserving class-relevant structure. Prior sampling rewards matching the encoded distribution to a simple distribution from which new latent codes can be drawn. A model that optimizes one of these objectives can easily fail another.

This tension is visible in variational auto-encoders (VAEs) [10], Wasserstein auto-encoders (WAEs) [14], and deep generative clustering methods such as VaDE [8]. VAEs provide a probabilistic latent-variable model, but their Euclidean Gaussian priors do not automatically form class-separated representations. WAEs match the aggregated posterior to a prior, often using MMD, but global matching alone does not impose semantic organization. VaDE introduces a Gaussian mixture prior for clustering, yet it remains Euclidean and does not explicitly separate class semantics from within-class style.

Hyperspherical latent variables offer a natural alternative for semantic organization. If class identity is encoded primarily by direction, then points on a unit sphere can be compared angularly and organized around class centers. Hyperspherical VAEs [4] and Spherical Cauchy VAEs [13] show that distributions on the sphere can be used in deep generative models. The single-latent CS-WAE model in this repository uses this idea: one spherical latent variable is matched to class-conditional Spherical Cauchy priors. It generates relatively clean MNIST samples, but its clustering performance is limited compared with the factorized model studied here.

We introduce F-CS-WAE, a factorized class-structured Spherical Cauchy WAE. The model decomposes the latent representation into two parts. The semantic latent z_c lives on a hypersphere and is aligned with class-conditional Spherical Cauchy priors. The style latent z_s lives in Euclidean space and is regularized toward a Gaussian prior. The architecture is motivated by a simple hypothesis: class identity should be represented by hyperspherical directions, while residual appearance variation should be carried by a separate style code.

The empirical picture is nuanced. On CIFAR-10, the factorized model dramatically improves clustering over available baselines and maintains substantially better sample quality than the older single-latent ResNet variant. On MNIST, it

obtains almost perfect clustering and excellent reconstructions, but naive prior samples are worse than those of CS-WAE. A targeted diagnostic shows why: z_s is globally matched to a Gaussian, but remains conditionally dependent on class. Thus, sampling z_c from a class prior and z_s independently from $N(0, I)$ creates latent pairs that the decoder did not learn to map to realistic images.

This paper is therefore written with a conservative claim. F-CS-WAE is a strong factorized model for label-guided clustering, and its failure mode is scientifically useful: it identifies conditional style mismatch as a central obstacle for making factorized WAE models simultaneously cluster-friendly and sample-friendly. For a final AAAI-27 submission, we argue that the correct target is not merely higher ACC/NMI/ARI, but a factorized objective that preserves clustering gains while repairing conditional decoding.

Contributions.

- We formulate F-CS-WAE, a factorized class-structured WAE with a Spherical Cauchy semantic latent and a Gaussian style latent.
- We introduce a four-phase training schedule that separates reconstruction learning, style regularization, semantic supervision, and class-prior alignment.
- We provide repository-local empirical evidence on CIFAR-10 and MNIST, including multi-seed CIFAR-10 results and a direct comparison against a single-latent CS-WAE variant.
- We diagnose a conditional prior mismatch in the MNIST model, showing that global style MMD is insufficient for reliable class-conditional generation.
- We outline the method extension required for an AAAI-27-strength paper: conditional style priors and per-class style matching.

2 Position in the Field

Auto-encoding generative models. VAEs [10] optimize an evidence lower bound with a KL penalty between approximate posteriors and a latent prior. WAEs [14] instead minimize reconstruction loss plus a discrepancy between the aggregated posterior and a target prior. This makes WAEs attractive for generative representation learning: the regularizer explicitly encourages encoded samples to be drawn from a distribution that is easy to sample. MMD-based WAEs use kernel two-sample distances [6] as the discrepancy.

Generative clustering. Deep Embedded Clustering [15] and VaDE [8] connect latent-variable modeling with clustering. VaDE in particular uses a Gaussian mixture prior, making cluster assignment part of the generative model. F-CS-WAE differs in two respects. First, it is label-guided rather than fully unsupervised: class labels are used during training to structure the semantic prior. Second, it uses hyperspherical semantic geometry rather than a Euclidean mixture.

Hyperspherical and heavy-tailed latent variables. Hyperspherical VAEs [4] motivate distributions on $\mathbb{S}^{d-1}$ when directional geometry is important. Spherical sliced-Wasserstein distances [1] adapt distribution matching ideas to spheres. Spherical Cauchy VAEs [13] further provide an efficient Mobius-transform reparameterization. F-CS-WAE uses this family for the semantic latent, but combines it with WAE-style MMD and label-guided class priors.

Factorization and disentanglement. Factorized latent spaces have a long history in disentanglement. InfoGAN [3] uses mutual information to make selected latent codes interpretable. FactorVAE [9] and β -TCVAE [2] penalize total correlation to encourage independent factors. These works motivate our central diagnostic: if semantic and style variables are intended to factorize, then matching only the marginal style distribution is insufficient. One must also control dependence between style and class or between style and semantic latents.

Where F-CS-WAE fits. F-CS-WAE sits between generative clustering and disentangled representation learning. Its semantic latent is explicitly class-structured, like a supervised version of mixture-prior clustering. Its style latent is intended to capture residual variation. Unlike purely disentangled VAEs, however, F-CS-WAE is evaluated not by interpretability alone but by clustering, reconstruction, and prior sampling. This combined evaluation is important because it exposes a failure mode that a clustering-only table would hide.

3 Method

3.1 Problem Setting

Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ be a labeled image dataset with K classes. The goal is to learn a generative representation that supports:

1. reconstruction: $x \mapsto \hat{x}$ should preserve visual content;
2. clustering: a simple clustering method on semantic codes should recover class structure;
3. sampling: latent samples from the prior should decode into realistic class-conditional images.

We call the setting *label-guided generative clustering*. Labels are used during training to organize the prior, but evaluation of clustering uses K-means and Hungarian matching on latent codes.

3.2 Architecture

F-CS-WAE uses a shared convolutional encoder with two heads:

$$E_\phi(x) = (\tilde{\mu}_c, r_c, \mu_s, \log \sigma_s^2). \quad (1)$$

The semantic direction is normalized,

$$\mu_c = \frac{\tilde{\mu}_c}{\|\tilde{\mu}_c\|_2}, \quad \rho_c = \sigma(r_c)(1 - \epsilon), \quad (2)$$

where $\mu_c \in \mathbb{S}^{d_c-1}$ and $\rho_c \in (0, 1)$. A semantic latent is sampled by a Spherical Cauchy Mobius reparameterization,

$$\eta_c \sim \text{Unif}(\mathbb{S}^{d_c-1}), \quad z_c = T_{\mu_c, \rho_c}(\eta_c). \quad (3)$$

The style latent uses a Gaussian reparameterization,

$$\epsilon_s \sim N(0, I), \quad z_s = \mu_s + \sigma_s \odot \epsilon_s. \quad (4)$$

The decoder receives the concatenated latent:

$$\hat{x} = D_\theta([z_c, z_s]). \quad (5)$$

In the current implementation, $d_c = 64$ and $d_s = 128$. The encoder uses a ResNet-18 body, and the decoder is a residual upsampling network. The style log-variance head is initialized near zero and clamped during sampling to avoid numerical overflow.

3.3 Class-Conditional Semantic Prior

For each class k , F-CS-WAE maintains one or more semantic centers

$$m_{k,r} \in \mathbb{S}^{d_c-1}, \quad k = 1, \dots, K, \quad r = 1, \dots, R. \quad (6)$$

The experiments in this draft use $R = 1$. A class prior sample is generated by selecting a center, drawing a uniform spherical noise vector, and applying the same Mobius transform:

$$\eta \sim \text{Unif}(\mathbb{S}^{d_c-1}), \quad z_{c,k}^p = T_{m_k, \rho_p}(\eta). \quad (7)$$

The prior radius is fixed at $\rho_p = 0.7$. Centers are updated by exponential moving average from encoded class means, which stabilizes class anchors while avoiding direct gradient competition with the decoder.

Algorithm 1 F-CS-WAE training loop

Require: Mini-batch (x, y) , encoder E_ϕ , decoder D_θ , class centers m_k

- 1: Encode $(\mu_c, \rho_c, \mu_s, \log \sigma_s^2) \leftarrow E_\phi(x)$
- 2: Sample $z_c \leftarrow T_{\mu_c, \rho_c}(\eta_c)$ with $\eta_c \sim \text{Unif}(\mathbb{S}^{d_c-1})$
- 3: Sample $z_s \leftarrow \mu_s + \sigma_s \odot \epsilon_s$ with $\epsilon_s \sim N(0, I)$
- 4: Decode $\hat{x} \leftarrow D_\theta([z_c, z_s])$
- 5: Sample class prior semantic latents $z_{c,k}^p$ from Spherical Cauchy centers m_k
- 6: Compute $\mathcal{L}_{\text{rec}}, \mathcal{L}_{\text{class}}, \mathcal{L}_{\text{agg}}, \mathcal{L}_{\text{style}}, \mathcal{L}_{\text{cls}}$
- 7: Apply phase weights $\alpha(t), \beta(t), \gamma(t), \eta(t)$ and update ϕ, θ, h
- 8: Accumulate μ_c by class and update centers m_k with EMA after each epoch

3.4 Loss Terms

The reconstruction objective is

$$\mathcal{L}_{\text{rec}} = \lambda_1 \|x - \hat{x}\|_1 + \lambda_{\text{LPIPS}} \text{LPIPS}(x, \hat{x}), \quad (8)$$

with $\lambda_1 = 1$ and $\lambda_{\text{LPIPS}} = 0.1$.

Semantic alignment uses a supervised class MMD:

$$\mathcal{L}_{\text{class}} = \frac{1}{K} \sum_{k=1}^K \text{MMD} \left(\{z_{c,i} : y_i = k\}, \{z_{c,k,j}^p\}_{j=1}^{n_k} \right). \quad (9)$$

An aggregated semantic MMD matches all semantic samples to a mixture of class priors:

$$\mathcal{L}_{\text{agg}} = \text{MMD} \left(\{z_{c,i}\}_{i=1}^b, \{z_{c,r_j,j}^p : r_j \sim \text{Unif}(1, \dots, K)\}_{j=1}^b \right). \quad (10)$$

Style matching uses a multi-scale Euclidean RBF MMD:

$$\mathcal{L}_{\text{style}} = \text{MMD}_{\text{RBF}} \left(\{z_{s,i}\}_{i=1}^b, \{\epsilon_i : \epsilon_i \sim N(0, I)\}_{i=1}^b \right). \quad (11)$$

An auxiliary classifier on μ_c gives

$$\mathcal{L}_{\text{cls}} = \text{CE}(h(\mu_c), y). \quad (12)$$

The total objective is

$$\mathcal{L} = \mathcal{L}_{\text{rec}} + \alpha(t)\mathcal{L}_{\text{class}} + \beta(t)\mathcal{L}_{\text{agg}} + \gamma(t)\mathcal{L}_{\text{style}} + \eta(t)\mathcal{L}_{\text{cls}}. \quad (13)$$

3.5 Four-Phase Training Schedule

Training proceeds in four phases:

- Phase A, epochs 0–49: reconstruction only.
- Phase B, epochs 50–99: add style MMD and auxiliary classification.
- Phase C, epochs 100–199: add weak semantic class and aggregate MMD.
- Phase D, epochs 200–299: full objective.

The schedule is designed to prevent early regularization from destroying reconstruction. The model first learns a useful decoder, then gradually shapes the semantic and style latents into samplable priors.

3.6 Sampling

For class-conditional generation, the current model samples

$$z_c \sim p(z_c \mid y = k), \quad z_s \sim N(0, I), \quad \hat{x} = D_\theta([z_c, z_s]). \quad (14)$$

This is the intended factorized generative story. The diagnostic in Section 5 shows that it is not fully achieved on MNIST: z_s remains class-dependent. This finding motivates the conditional style extension in Section 6.

Table 1: CIFAR-10 F-CS-WAE results over three seeds. ACC, NMI, ARI, SSIM are shown as fractions; lower FID and LPIPS are better.

Seed	ACC	NMI	ARI	FID	LPIPS	SSIM	PSNR
0	0.8043	0.6426	0.6259	83.83	0.2126	0.7073	21.88
1	0.8160	0.6607	0.6452	83.45	0.2183	0.6967	21.67
2	0.8059	0.6455	0.6264	81.85	0.2137	0.7060	21.76
Mean	0.8087	0.6496	0.6325	83.04	0.2149	0.7034	21.77
Std	0.0052	0.0079	0.0090	0.86	0.0025	0.0047	0.09




Figure 1: CIFAR-10 clustering and generation metrics across the three completed F-CS-WAE seeds.

4 Experiments

4.1 Datasets and Evaluation

We report completed repository-local experiments on CIFAR-10 [11] and MNIST [12]. CIFAR-10 is the primary completed multi-seed setting for F-CS-WAE; MNIST is used for detailed analysis because it exposes the style-prior failure mode clearly.

Clustering uses K-means on the semantic means μ_c , followed by Hungarian matching to labels. We report ACC, NMI, and ARI. Generation quality is measured by FID [7]; reconstruction quality uses SSIM, PSNR, and LPIPS [16]. Unless otherwise stated, FID uses 10,000 generated and 10,000 real images.

4.2 CIFAR-10 Main Result

Table 1 reports the completed three-seed CIFAR-10 results. F-CS-WAE is stable across seeds: clustering accuracy varies by only 0.52 percentage points, and FID by less than one point.

4.3 CIFAR-10 Comparison to Available Baselines

Table 2 compares F-CS-WAE to available local baselines. Because not all baselines have three-seed runs, the table explicitly reports the source status. The comparison should be treated as evidence for the current repository protocol rather than as a final benchmark table.

The gap is large: F-CS-WAE improves ACC by roughly 59 percentage points over VAE/WAE-MMD and by roughly 66 points over VaDE in this protocol. The FID gap is also favorable relative to these baselines. This is the strongest current evidence for F-CS-WAE as a paper direction.

Table 2: Available CIFAR-10 comparison. F-CS-WAE is a 3-seed mean; other rows are completed local seed-0 results unless marked otherwise.

Method	Status	ACC $\uparrow$	NMI $\uparrow$	ARI $\uparrow$	FID $\downarrow$
F-CS-WAE	3-seed mean	0.8087	0.6496	0.6325	83.04
F-CS-WAE full ablation run	completed seed 0	0.8149	0.6622	0.6438	83.59
ResNetAE	seed 0	0.2202	0.1017	0.0493	121.01
VAE	seed 0	0.2220	0.0955	0.0511	179.98
WAE-MMD	seed 0	0.2099	0.0979	0.0469	235.43
VaDE	seed 0	0.1475	0.0427	0.0151	475.87




Figure 2: Training curves for the completed CIFAR-10 F-CS-WAE runs. The four-phase schedule keeps training stable after style and semantic regularizers are introduced.

4.4 Qualitative CIFAR-10 Structure

Figure 3 shows class-conditional prior samples from the completed CIFAR-10 checkpoint. Figure 4 visualizes the intended factorization: rows with fixed semantics and varying styles, and rows with fixed style and varying semantic centers. These figures are not a substitute for quantitative ablations, but they help inspect whether the decoder uses z_c and z_s in a plausible way.

4.5 MNIST: High Clustering, Poor Naive Prior Sampling

MNIST exposes the central trade-off more sharply. Table 3 compares F-CS-WAE seed 0 to the single-latent CS-WAE seed 0. F-CS-WAE strongly improves clustering and reconstruction metrics, but its FID is worse. This is the key warning for an AAAI-27 submission: ACC/NMI/ARI alone would overstate the generative quality of the factorized model.

5 Failure Analysis: Conditional Style Mismatch

The MNIST result raises a question: why does a model with excellent reconstruction and clustering produce poor prior samples? We tested whether the style latent is globally matched to $N(0, I)$ and whether it is independent of class. The answer is mixed.

On 2,048 MNIST test samples encoded by the F-CS-WAE seed-0 checkpoint, the posterior style samples have global standard deviation approximately 0.989, per-dimension variance approximately 0.973, and multi-scale MMD to $N(0, I)$ approximately 0.0017. By a marginal aggregate test, style matching appears successful. However, class-conditional style means remain far apart: the average pairwise distance between class style means is approximately 6.25, while within-class style norm is approximately 10.25. In other words, z_s still carries class information.

This explains the generation failure. During training, the decoder sees pairs (z_c, z_s) from the same image. If z_s contains class information, then the decoder learns class-specific interactions between z_c and z_s . During generation, the current




Figure 3: CIFAR-10 class-conditional samples from F-CS-WAE priors. Each row corresponds to a class prior.

prior samples z_c from class k but samples z_s independently from the global Gaussian. This creates combinations that may never occur under the encoded posterior.

Figure 6 visualizes the issue. Empirical same-class style codes decode cleanly, while independent Gaussian styles often create broken samples. Setting $z_s = 0$ also fails to preserve class semantics. The conclusion is direct: current F-CS-WAE improves semantic clustering, but its intended factorization is incomplete because the style latent is not class-invariant.

6 Toward an AAI-27-Strength Method

The current draft supports the following claim: factorized semantic/style modeling can greatly improve clustering, but naive marginal style matching is insufficient for reliable class-conditional generation. For an AAI-27 submission, the method should close this gap. We propose three concrete extensions.

6.1 Conditional Style Prior

Replace the independent prior $z_s \sim N(0, I)$ with a class-conditional style prior:

$$z_s \sim p(z_s \mid y = k). \quad (15)$$

Table 3: MNIST seed-0 comparison. CS-WAE is the single-latent Spherical Cauchy WAE variant; F-CS-WAE is the factorized model.

Method	ACC	NMI	ARI	FID	LPIPS	SSIM	PSNR
CS-WAE	0.8735	0.8863	0.8411	26.95	0.0603	0.8127	16.38
F-CS-WAE	0.9915	0.9758	0.9813	73.99	0.0120	0.9833	26.88

The simplest version uses an EMA Gaussian per class:

$$p(z_s | y = k) = N(\bar{\mu}_{s,k}, \text{diag}(\bar{\sigma}_{s,k}^2)). \quad (16)$$

This extension matches the empirical finding that style is currently class-dependent. It is likely to improve samples quickly, but it changes the interpretation: style is no longer strictly class-independent. The paper can still make a strong claim if it frames the model as a *conditional factorized prior* rather than a fully disentangled semantic/style model.

6.2 Per-Class Style MMD

If the goal is class-invariant style, then the loss should align each conditional style aggregate to the same Gaussian:

$$\mathcal{L}_{\text{style-class}} = \frac{1}{K} \sum_{k=1}^K \text{MMD}(\{z_{s,i} : y_i = k\}, \{\epsilon_j : \epsilon_j \sim N(0, I)\}). \quad (17)$$

This directly targets the observed failure. It is stronger than global style MMD because it prevents class means from separating in style space while still matching the aggregate.

6.3 Style Adversary or HSIC Penalty

A complementary approach is to penalize dependence between z_s and class labels. A gradient-reversal classifier can try to predict y from z_s , while the encoder tries to make this prediction fail. Alternatively, an HSIC-style dependence penalty [5] can reduce dependence between z_s and labels or between z_s and z_c .

6.4 What the Final AAI Claim Should Be

A paper based only on ACC/NMI/ARI would be vulnerable because the model is explicitly generative. The final AAI-level claim should be:

Factorized Spherical Cauchy WAE models improve label-guided clustering, but require conditional style alignment to preserve prior sample quality; with this alignment, the model achieves strong clustering and competitive generation under a unified class-conditional generative objective.

This claim is stronger than a pure clustering result because it contributes both a model and a diagnosis of why factorization can fail.

7 Discussion

Why CIFAR-10 looks stronger than MNIST. On CIFAR-10, F-CS-WAE improves both clustering and FID relative to the available local baselines. This likely reflects the weakness of the baselines under the same ResNet/protocol and the benefit of having a high-capacity factorized representation. On MNIST, the single-latent CS-WAE prior is already well matched to the simple digit manifold, so the factorized model must justify its extra degrees of freedom. It improves clustering, but the style latent becomes an unregulated route for class information.

Why reconstruction is misleading. The MNIST F-CS-WAE checkpoint has excellent SSIM, PSNR, and LPIPS, yet poor prior samples. This confirms a standard lesson in generative modeling: reconstruction quality is not equivalent to sample quality. A decoder can reconstruct encoded latents that lie on the posterior manifold while failing on samples from the intended prior.

Why the failure is useful. The failure is not merely a bad hyperparameter setting. It identifies a missing conditional constraint in the factorized objective. Global matching $q(z_s) \approx N(0, I)$ does not imply $q(z_s | y = k) \approx N(0, I)$ for all k , nor does it imply $z_s \perp y$. This is precisely the gap that conditional style priors or per-class style MMD should close.

Limitations of the current draft. The completed evidence is not yet sufficient for a final AAAI-27 submission. CIFAR-10 has strong three-seed results, but MNIST and Fashion-MNIST need full multi-seed factorized runs. Ablations are incomplete. The conditional style fix proposed in Section 6 must be implemented and evaluated. The current paper draft is therefore a detailed research manuscript and submission plan, not a finished empirical package.

8 Conclusion

We presented F-CS-WAE, a factorized class-structured Spherical Cauchy WAE for label-guided generative clustering. The model separates semantic and style latents, aligns semantics with hyperspherical class priors, and regularizes style toward a Gaussian prior through MMD. Completed CIFAR-10 experiments show strong and stable clustering performance with favorable generation metrics relative to available baselines. MNIST diagnostics reveal a deeper issue: marginal style matching can hide class dependence in the style latent, causing naive independent prior samples to decode poorly.

The main lesson is that factorization is not obtained by architectural separation alone. A factorized generative model needs conditional distributional constraints that match how sampling will be performed. For an AAAI-27-strength submission, the next step is therefore clear: preserve the clustering gains of F-CS-WAE while adding conditional style alignment to make class-conditional generation reliable.

References

- [1] Clément Bonet, Paul Berg, Nicolas Courty, François Septier, Lucas Drumetz, and Minh-Tan Pham. Spherical sliced-wasserstein. *arXiv preprint arXiv:2206.08780*, 2022.
- [2] Ricky T. Q. Chen, Xuechen Li, Roger Grosse, and David Duvenaud. Isolating sources of disentanglement in variational autoencoders. *arXiv preprint arXiv:1802.04942*, 2018.
- [3] Xi Chen, Yan Duan, Rein Houthoofd, John Schulman, Ilya Sutskever, and Pieter Abbeel. InfoGAN: Interpretable representation learning by information maximizing generative adversarial nets. *arXiv preprint arXiv:1606.03657*, 2016.
- [4] Tim R. Davidson, Luca Falorsi, Nicola De Cao, Thomas Kipf, and Jakub M. Tomczak. Hyperspherical variational auto-encoders. *arXiv preprint arXiv:1804.00891*, 2018.
- [5] Arthur Gretton, Olivier Bousquet, Alex Smola, and Bernhard Schölkopf. Measuring statistical dependence with hilbert-schmidt norms. In *International Conference on Algorithmic Learning Theory*, pages 63–77. Springer, 2005.
- [6] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola. A kernel two-sample test. *Journal of Machine Learning Research*, 13:723–773, 2012.
- [7] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. GANs trained by a two time-scale update rule converge to a local nash equilibrium. In *Advances in Neural Information Processing Systems*, 2017.
- [8] Zhuxi Jiang, Yin Zheng, Huachun Tan, Bangsheng Tang, and Hanning Zhou. Variational deep embedding: An unsupervised and generative approach to clustering. *arXiv preprint arXiv:1611.05148*, 2016.
- [9] Hyunjik Kim and Andriy Mnih. Disentangling by factorising. *arXiv preprint arXiv:1802.05983*, 2018.
- [10] Diederik P. Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [11] Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- [12] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.

- [13] Lukas Sablica and Kurt Hornik. Hyperspherical variational autoencoders using efficient spherical cauchy distribution. *arXiv preprint arXiv:2506.21278*, 2025.
- [14] Ilya Tolstikhin, Olivier Bousquet, Sylvain Gelly, and Bernhard Schölkopf. Wasserstein auto-encoders. *arXiv preprint arXiv:1711.01558*, 2017.
- [15] Junyuan Xie, Ross Girshick, and Ali Farhadi. Unsupervised deep embedding for clustering analysis. *arXiv preprint arXiv:1511.06335*, 2015.
- [16] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018.

A Repository Artifacts Used

All quantitative claims in this draft are taken from local repository artifacts:

- CIFAR-10 F-CS-WAE metrics: `runs_f/cifar10/seed_{0,1,2}/metrics.json`.
- CIFAR-10 summary tables and figures: `paper_outputs/f_cs_wae_cifar10/`.
- MNIST F-CS-WAE metrics and images: `runs_f/mnist/seed_0/`.
- MNIST CS-WAE metrics and images: `runs/mnist/seed_0/`.

B AAI-27 Submission Checklist

Before submission, the following experiments should be completed:

- Run F-CS-WAE on MNIST and Fashion-MNIST for seeds 0, 1, 2.
- Implement conditional style prior and per-class style MMD.
- Run ablations: no style latent, no semantic MMD, no style MMD, no auxiliary classifier, conditional style prior, per-class style MMD, style adversary, style dimension sweep.
- Recompute all FID directories after changing sampling.
- Add a final benchmark table with matched seeds for all baselines.
- Replace this preamble with the official AAI-27 author kit.

Style-Semantic Diagnostic: fixed z_c varies style; fixed z_s varies class center





Figure 4: CIFAR-10 style/semantic diagnostic. The figure probes whether varying z_s at fixed z_c , and varying z_c at fixed z_s , produces interpretable changes.



(a) CS-WAE prior samples.



(b) F-CS-WAE naive independent prior samples.

Figure 5: MNIST generation comparison. The factorized model clusters and reconstructs better, but independent sampling of $z_s \sim N(0, I)$ creates unrealistic or class-mismatched samples.



Figure 6: MNIST style-prior diagnostic. For each digit, three rows compare $z_s \sim N(0, I)$, $z_s = 0$, and empirical posterior z_s sampled from the same class. Empirical same-class style codes recover clean digits; independent Gaussian style codes often decode as broken or wrong-class samples.